

Deep Learning


WaveNet: From article to TensorFlow code

Tero Keski-Valkama



2010-01-19

- WaveNet: A Generative Model for Raw Audio by Aaron van den Oord et al. @ DeepMind: <https://arxiv.org/pdf/1609.03499.pdf>
- WaveNet Tensorflow implementation by Igor Babuschkin et al: <https://github.com/ibab/tensorflow-wavenet>



wavenet/model.py#L172

```
def _create_causal_layer(self, input_batch):  
    '''Creates a single causal convolution layer.  
    The layer can change the number of channels.  
    '''  
    with tf.name_scope('causal_layer'):  
        weights_filter = self.variables['causal_layer']['filter']  
    return causal_conv(input_batch, weights_filter, 1)
```

Figure 4: Overview of the residual block and the entire architecture.

wavenet/ops.py#L27




Figure 2: Visualization of a stack of causal convolutional layers.





Figure 3: Visualization of a stack of *dilated* causal convolutional layers.

```
def time_to_batch(value, dilation, name=None):
    with tf.name_scope('time_to_batch'):
        shape = tf.shape(value)
        pad_elements = dilation - 1 - (shape[1] + dilation - 1) % dilation
        padded = tf.pad(value, [[0, 0], [0, pad_elements], [0, 0]])
        reshaped = tf.reshape(padded, [-1, dilation, shape[2]])
        transposed = tf.transpose(reshaped, perm=[1, 0, 2])
        return tf.reshape(transposed, [shape[0] * dilation, -1, shape[2]])

def batch_to_time(value, dilation, name=None):
    with tf.name_scope('batch_to_time'):
        shape = tf.shape(value)
        prepared = tf.reshape(value, [dilation, -1, shape[2]])
        transposed = tf.transpose(prepared, perm=[1, 0, 2])
        return tf.reshape(transposed,
                           [tf.div(shape[0], dilation), -1, shape[2]])

def causal_conv(value, filter_, dilation, name='causal_conv'):
    with tf.name_scope(name):
        # Pad beforehand to preserve causality.
        filter_width = tf.shape(filter_)[0]
        padding = [[0, 0], [(filter_width - 1) * dilation, 0], [0, 0]]
        padded = tf.pad(value, padding)
        if dilation > 1:
            transformed = time_to_batch(padded, dilation)
            conv = tf.nn.conv1d(transformed, filter_, stride=1, padding='SAME')
            restored = batch_to_time(conv, dilation)
        else:
            restored = tf.nn.conv1d(padded, filter_, stride=1, padding='SAME')
        # Remove excess elements at the end.
        result = tf.slice(restored,
                           [0, 0, 0],
                           [-1, tf.shape(value)[1], -1])

    return result
```



wavenet/model.py#L300

```
# Add all defined dilation layers.
with tf.name_scope('dilated_stack'):
    for layer_index, dilation in enumerate(self.dilations):
        with tf.name_scope('layer{}'.format(layer_index)):
            output, current_layer = self._create_dilation_layer(
                current_layer, layer_index, dilation)
            outputs.append(output)
```

Figure 4: Overview of the residual block and the entire architecture.

wavenet/model.py#L181




Figure 4: Overview of the residual block and the entire architecture.

```
def _create_dilation_layer(self, input_batch, layer_index, dilation):
    """Creates a single causal dilated convolution layer.
    The layer contains a gated filter that connects to dense output
    and to a skip connection:
        |-> [gate]   -|           |-> 1x1 conv -> skip output
        |           |-> (*) -|
    input -|-> [filter] -|           |-> 1x1 conv -|
        |           |-----|           |-> (+) -> dense output
        |-----|

    Where '[gate]' and '[filter]' are causal convolutions with a
    non-linear activation at the output.
    """
    variables = self.variables['dilated_stack'][layer_index]
    weights_filter = variables['filter']
    weights_gate = variables['gate']
    conv_filter = causal_conv(input_batch, weights_filter, dilation)
    conv_gate = causal_conv(input_batch, weights_gate, dilation)
    ... CODE FOR OPTIONAL BIASES ...
    out = tf.tanh(conv_filter) * tf.sigmoid(conv_gate)
    # The 1x1 conv to produce the residual output
    weights_dense = variables['dense']
    transformed = tf.nn.conv1d(
        out, weights_dense, stride=1, padding="SAME", name="dense")
    # The 1x1 conv to produce the skip output
    weights_skip = variables['skip']
    skip_contribution = tf.nn.conv1d(
        out, weights_skip, stride=1, padding="SAME", name="skip")
    ... CODE FOR OPTIONAL BIASES / HISTOGRAM OUTPUT ...
    return skip_contribution, input_batch + transformed
```




Figure 4: Overview of the residual block and the entire architecture.

wavenet/model.py#L127

```
current['dense'] = create_variable(  
    'dense',  
    [1,  
     self.dilation_channels,  
     self.residual_channels])  
current['skip'] = create_variable(  
    'skip',  
    [1,  
     self.dilation_channels,  
     self.skip_channels])
```

wavenet/model.py#L212

```
# The 1x1 conv to produce the residual output  
weights_dense = variables['dense']  
transformed = tf.nn.conv1d(  
    out, weights_dense, stride=1, padding="SAME", name="dense")  
# The 1x1 conv to produce the skip output  
weights_skip = variables['skip']  
skip_contribution = tf.nn.conv1d(  
    out, weights_skip, stride=1, padding="SAME", name="skip")
```




Figure 4: Overview of the residual block and the entire architecture.

wavenet/model.py#L155

```
current['postprocess1'] = create_variable(  
    'postprocess1',  
    [1, self.skip_channels, self.skip_channels])  
current['postprocess2'] = create_variable(  
    'postprocess2',  
    [1, self.skip_channels, self.quantization_channels])
```

wavenet/model.py#L324

```
w1 = self.variables['postprocessing']['postprocess1']  
w2 = self.variables['postprocessing']['postprocess2']  
... CODE FOR OPTIONAL BIASES / HISTOGRAM OUTPUT ...  
# We skip connections from the outputs of each layer, adding them  
# all up here.  
total = sum(outputs)  
transformed1 = tf.nn.relu(total)  
conv1 = tf.nn.conv1d(transformed1, w1, stride=1, padding="SAME")  
... CODE FOR OPTIONAL BIASES ...  
transformed2 = tf.nn.relu(conv1)  
conv2 = tf.nn.conv1d(transformed2, w2, stride=1, padding="SAME")
```